_____
University of Texas at Dallas
CS 6322 : Information Retrieval
Fall 2014
Homework # 2

Instructor: Dr. Sanda Harabagiu
Grader: Travis Goodwin

Issued: October 6th 2014
Due October 27th 2014 before midnight
_____

Problem (200 points)
Index building

_____

   In this assignment you build two versions of the index for a simple statistical retrieval system and also each version of the index shall be in uncompressed form and compressed form. In the next assignment, you will build the retrieval system itself. For this assignment, index the Cranfield documents used in the last assignment:
A copy of the publicly available Cranfield collection is located on the UTD Apache machine at:
   /people/cs/s/sanda/cs6322/Cranfield/

   You need to build: Index_Version1.uncompress and Index_Version2.uncompress of the index.
   Version 1 of your index considers the terms in the dictionary to me lemmas of words, whereas version 2 of your index considers that the terms of the dictionary are stems of the tokens you have processed in Homework 1.

Do not store stop-words in any version of your index. Before building the dictionaries of any of the two versions of index, it is recommended to remove the stop words. You may use the stop-word list from the directory on the UTD Apache machine at:
   /people/cs/s/sanda/cs6322/Cranfield/resourcesIR.

The terms in your 2nd version of the index should be stemmed with the Porter stemmer. The code for the Porter stemmer is available in the same directory on the UTD Apache machine at:
   /people/cs/s/sanda/cs6322/resourcesIR.

For every term that is both versions of the indexed, store:
   - Document frequency (df): The number of documents that the term occurs in.
   - Term frequency (tf): The number of times that the term occurs in
        each document, and
   - The list of documents containing the term.

For each document, store the frequency of the most frequent stem in that document (max_tf), and the total number of word occurrences in the document (doclen). To be noted that the value of doclen includes the number of stop-words encountered in the respective documents.

Store the inverted lists in your own storage manager.

You also are required to build the compressed versions of your index: Index_Version1.compressed and Index_Version2.compressed. To do so, you shall (a) compress the dictionaries of both versions of the index and compress the inverted lists before storing them, using delta encoding for the document-id and gamma code for the frequency information. You can compress the dictionary with any method that we have presented in class.

A penalty of -100 points will be applied if you do not have also the compressed versions of the two forms of index. If you return only one version of an uncompressed index, you shall obtain a maximum of 50 points in this homework, if you return only one version, both in compress and uncompressed form, you will obtain a maximum of 100 points in this homework.

Delta codes are similar to the gamma codes: they represent a gap by a pair: (length, offset). First the number is represented in binary code. The length of the binary representation is encoded in gamma code, prior to removing the leading 1-bit. After generating the code of the length only, the leading 1-bit is removed and represented in gamma code.

Example 1: To write 5 in gamma and delta codes we perform the following operations:
1. write 5 in binary as 101
2. For the gamma code remove the leading 1-bit to obtain the offset: 01
3. The length of the offset is 2:
   - In unary the length is 110
4. The code of 5 in gamma is 11001 (or 110,01 – to represent [length, offset])
5. For the delta code, the length of the offset is 3, because the leading 1-bit is removed afterwards. When writing the length=3 in gamma code it becomes:10,1
6. for delta code, the leading 1-bit of the offset is removed now, generating an offset of 10
7. The code of 5 in delta is 10101

Example 2: To write 9 in gamma and delta codes we perform the following operations:
1. write 9 in binary as 1001
2. for gamma code, remove the leading 1-bit to obtain the offset: 001
3. The length of the offset is 3:
   - In unary the length is 1110
4. The code of 9 in gamma is 1110001 (or 1110,001 – to represent [length, offset])
5. for the delta code, the length of the binary representation is 4
6. The length is represented in gamma code: 11000

7. The leading 1-bit of the binary representation is removed, generating the offset 001
8. The code of 9 in delta is 11000001

Example 3: To write 1 in gamma and delta , we perform the following operations:
1. write 1 in binary: 1
2. For the gamma code, we remove 1, generating a length=0, which is still 0 in unary code
3. in Gamma code, 1 becomes 0
4. for the Delta code, the length of the binary representation for 1 is 1. Gamma code for 1 is 0
5. The code for 1 in delta is 0

More values for gamma and delta codes are given in the following table:

N	Binary	Gamma code	Delta code
1	1	Len(-)=0; unary(0)=0; Gamma(1)=0	Len(1)=1; Gamma(1)=0; Delta(1)=0
2	10	Len(0)=1; unary(1)=10; Gamma(2)=100	Len(10)=2; Gamma(2)=100; Delta(2)=1000
3	11	Len(1)=1; unary(1)=10; Gamma(3)=101	Len(11)=2; Gamma(2)=100; Delta(3)=1001
4	100	Len(00)=2; unary(2)=110; Gamma(4)=11000	Len(100)=3; Gamma(3)=101; Delta(4)=10100
5	101	Len(01)=2; unary(2)=110; Gamma(5)=11001	Len(101)=3; Gamma(3)=101; Delta(5)=10101
6	110	Len(10)=2; unary(2)=110; Gamma(6)=11010	Len(110)=3; Gamma(3)=101; Delta(6)=10110
7	111	Len(11)=2; unary(2)=110; Gamma(7)=11011	Len(111)=3; Gamma(3)=101; Delta(7)=10111
8	1000	Len(000)=3; unary(3)=1110; Gamma(8)=111000	Len(1000)=4; Gamma(4)=11000; Delta(8)=11000000
9	1001	Len(001)=3; unary(3)=1110; Gamma(9)=111001	Len(1001)=4; Gamma(4)=11000; Delta(9)=11000001
10	1010	Len(010)=3; unary(2)=1110; Gamma(10)=1110010	Len(1010)=4; Gamma(4)=11000; Delta(10)=11000010

Production-level IR systems build these compressed indices in about 5 minutes. If your program takes more than an hour, you are doing something wrong.

Turn in your program, a written description of your program, and the following statistics:
  - the elapsed time ("wall-clock time") required to build any version of your index,

- the size of the index Version 1 uncompressed (in bytes),
 - the size of the index Version 2 uncompressed (in bytes),
 - the size of the index Version 1 compressed (in bytes),
 - the size of the index Version 2 compressed (in bytes),
 - the number of inverted lists in each version of the index, and
 - the df, tf, and inverted list length (in bytes) for the terms:
   "Reynolds", "NASA", "Prandtl", "flow", "pressure", "boundary", "shock" (or stems
that correspond to them).